

Project 3

Table of Contents

Table of Contents	1
Introduction	1
Folders structure	1
Installation	2
Requirements	2
Installation	2
Compile	2
Usage	2
Test	2
🖥️ Hardware Specifications	2
Problem 1	3
🔧 Experiment Setup	3
📊 Performance Chart	4
📊 Detailed Observations	4
1. Execution Time (Lower is Better)	4
2. Performance Index (Higher is Better)	4
🧠 Explanation of Results	5
✅ Why Dynamic Scheduling + Chunk Size 10 Performs Best	5
⚠️ Why Performance Drops or Plateaus at 14–16 Threads	5
🏁 Conclusion	5

Introduction

This project was written in **C++** and uses **OpenMP**

Folders structure

```
|— xmake.lua
|— bin -> folder created after compilation
|   |— prob1 -> first problem executable
|   |— prob2 -> second problem executable
|— problem1 -> contains the code for the first problem
|   |— xmake.lua
|   |— prob1.cpp -> problem 1 source code
|— problem2 -> contains the code for the second problem
|   |— xmake.lua
|   |— prob2.cpp -> problem 2 source code
```

Installation

Requirements

- C++ compiler (g++, clang++, msvc++)
- XMake
- Git
- OpenMP

Installation

```
git clone git@github.com:GlassAlo/CAU_Multicore.git
cd CAU_Multicore/proj3
```

Compile

```
xmake f -m release && xmake -y
```

- **-m release** is used to compile the project in release mode, which is faster than debug mode.
- **xmake -y** is used to compile the project. The **-y** flag is used to skip the confirmation prompt.
- **xmake** will create a **bin** folder with the executables inside.

Usage

```
./bin/prob1 <scheduling type> <number of threads>
./bin/prob2 <number end>
```

Test

Hardware Specifications

- **OS:** Garuda Linux Broadwing x86_64
- **Kernel:** 6.13.8-zen1-1-zen
- **CPU:** AMD Ryzen 9 5900HS with Radeon Graphics (16) @ 4.680GHz
 - Cores 8
 - Uniform core design
 - Threads 16
 - Base clock 3.0GHz
 - Max boost clock up to 4.6GHz
 - L3 cache 16MB
 - Memory PCIe 3.0

- Supports Simultaneous Multithreading (SMT), with each cores supporting two threads
 - **Integrated GPU:** AMD ATI Radeon Vega Series / Radeon Vega Mobile Series
 - **Discrete GPU:** NVIDIA GeForce RTX 3080 Mobile / Max-Q (8GB/16GB)
 - **RAM:** 32GB
 - **Disk:** 1TB SSD
 - **Shell:** zsh
- Using Arch Linux comes with a cost, the CPU pilots might not be very efficient, stable or up to date.

Problem 1

Experiment Setup

We measured the execution time and derived performance index using:

- **Static load balancing with default chunk size**
- **Dynamicload balancing with default chunk size**
- **Static load balancing with a chunk size of 10**
- **Dynamic load balancing (task size 10)**

Each implementation was tested using 1, 2, 4, 6, 8, 10, 12, 14 and 16 threads, with execution times measured in milliseconds (ms).

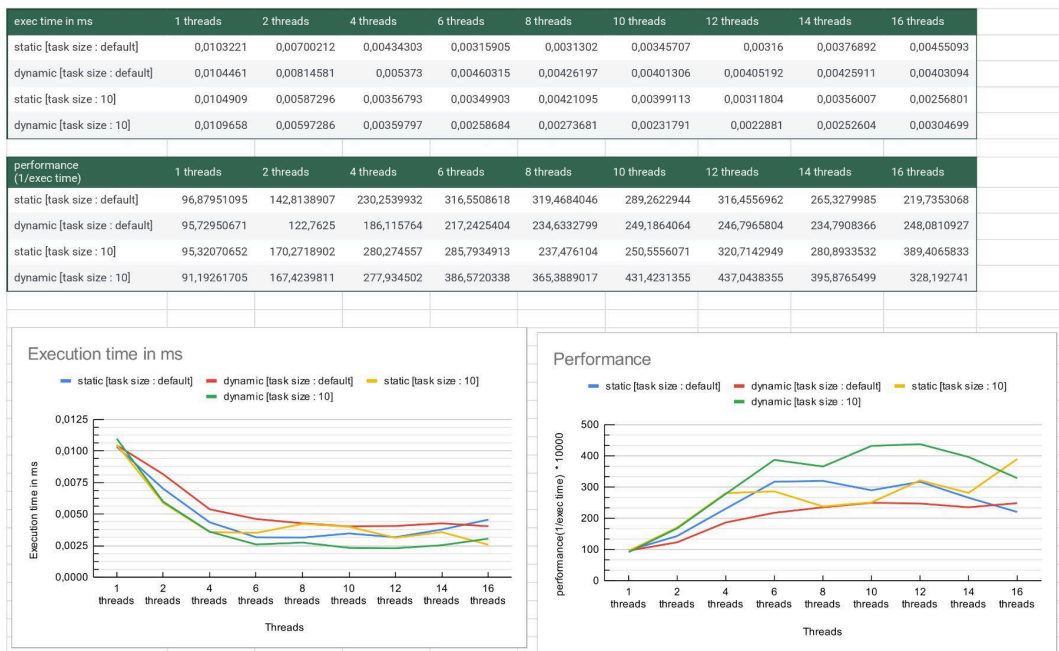
The performance index was defined as:

- **Performance Index = (1/Execution Time)**

This means that higher values indicate better performance.

- All the tests were done with the same number of threads and the same end number.
- The end number was set to 200000.

Performance Chart



Detailed Observations

1. Execution Time (Lower is Better)

- All four scheduling strategies show significant improvements from 1 to 8 threads.
- **Dynamic (chunk size = 10)** shows the fastest execution at most thread counts beyond 4.
- **Static (default chunk size)** performs well until 8–10 threads, after which gains taper off.
- After 12 threads, **all methods plateau or degrade**, showing signs of overhead or saturation.

2. Performance Index (Higher is Better)

- **Dynamic with chunk size 10** consistently achieves the highest performance from 4 to 12 threads.
- **Static with chunk size 10** also shows strong performance but with more fluctuation.
- The **default static** and **default dynamic** approaches underperform with high thread counts, likely due to inefficient load balancing or lack of granularity.

Explanation of Results

Why Dynamic Scheduling + Chunk Size 10 Performs Best

- **Better Load Balancing:** Tasks are dynamically pulled by idle threads, reducing idle time.
- **Fine Granularity:** Smaller chunks allow better utilization of threads, especially in unbalanced workloads like recursive tasks.

Why Performance Drops or Plateaus at 14–16 Threads

- **Overhead of too many tasks:** With 200,000 elements and chunk size 10, you create 20,000 tasks — overhead increases with thread count.
- **Diminishing returns:** CPU reaches saturation around 8–12 threads, especially with hyperthreading.
- **Contention and synchronization costs** (e.g., **reduction**, task queue contention) rise with more threads.

Conclusion

Dynamic scheduling with small chunk sizes (10) provides the best performance in highly parallel divide-and-conquer tasks on modern CPUs with SMT.

Static scheduling performs reasonably well but lacks adaptability to runtime imbalances. While increasing threads generally improves performance, **beyond 12 threads** there are **diminishing returns** due to overhead, memory bandwidth saturation, and synchronization costs.

The results validate that **task granularity and scheduling policy** are critical to achieving optimal parallel performance in OpenMP.