

# Problem1

## Table of Contents

<b>Table of Contents</b>	<b>1</b>
Introduction	1
Folders structure	1
Installation	2
Requirements	2
Installation	2
Compile	2
Usage	3
Test	3
 Hardware Specifications	3
Prime number checker	4
 Experiment Setup	4
 Performance Chart	4
 Interpretation	5
 Conclusion	5

## Introduction

This project was written in **C++** because it's a language that I enjoy using and I am familiar with. It includes an architecture made to be extensible and easy to maintain while reusing as much code as possible which means a bit more overhead.

## Folders structure

```
|— xmake.lua
|— bin -> folder created after compilation
|   |— MatmultD -> second problem executable
|   |— pc_dynamic -> first problem executable (dynamic)
|   |— pc_static_block -> first problem executable (static block)
|   |— pc_static_cyclic -> first problem executable (static cyclic)
|— xmake.lua
|— problem1 -> contains the code for the first problem
|   |— xmake.lua
|   |— pc_serial.java.cpp -> translation of the basic java code to cpp
|   |— dynamic
|   |   |— xmake.lua
|   |   |— Main.cpp
|   |   |— DynamicThread.cpp and DynamicThread.hpp -> Child class of Thread
|   |   |— WorkQueue.cpp and WorkQueue.hpp
```

```

|   ├── static_block
|   |   ├── xmake.lua
|   |   ├── Main.cpp
|   |   ├── StaticBlockThread.cpp and StaticBlockThread.hpp -> Child class of Thread
|   ├── static_cyclic
|   |   ├── xmake.lua
|   |   ├── Main.cpp
|   |   ├── StaticCyclicThread.cpp and StaticCyclicThread.hpp -> Child class of Thread
|   └── problem2 -> contains the code for the second problem
|       ├── xmake.lua
|       ├── Main.cpp
|       ├── Matrix.cpp and Matrix.hpp
|       ├── MatrixThread.cpp and MatrixThread.hpp -> Child class of Thread
|   └── Shared -> contains the code shared between the problems
|       ├── Clock.cpp and Clock.hpp
|       ├── PrimeChecker.cpp and PrimeChecker.hpp
|       ├── Thread.cpp and Thread.hpp
|       ├── ThreadPool.hpp -> No cpp file because it's a template class
|   └── data -> folder containing the test matrices

```

## Installation

### Requirements

- C++ compiler (g++, clang++, msvc++)
- XMake
- Git

### Installation

```

git clone git@github.com:GlassAlo/CAU_Multicore.git
cd CAU_Multicore/proj1

```

### Compile

```

xmake f -m release && xmake -y

```

- **-m release** is used to compile the project in release mode, which is faster than debug mode.
- **xmake -y** is used to compile the project. The **-y** flag is used to skip the confirmation prompt.
- **xmake** will create a **bin** folder with the executables inside.

## Usage

```
./bin/pc_dynamic <number of threads> <number end number>  
./bin/pc_static_block <number of threads> <number end number>  
./bin/pc_static_cyclic <number of threads> <number end number>
```

- **The first three executables are for the first problem**, which is a prime number checker.
- The first three executables take two arguments: the number of threads and the end number.

## Test



### Hardware Specifications

- **OS:** Garuda Linux Broadwing x86\_64
  - **Kernel:** 6.13.8-zen1-1-zen
  - **CPU:** AMD Ryzen 9 5900HS with Radeon Graphics (16) @ 4.680GHz
    - Cores 8
      - Uniform core design
    - Threads 16
    - Base clock 3.0GHz
    - Max boost clock up to 4.6GHz
    - L3 cache 16MB
    - Memory PCIe 3.0
    - Supports Simultaneous Multithreading (SMT), with each cores supporting two threads
  - **Integrated GPU:** AMD ATI Radeon Vega Series / Radeon Vega Mobile Series
  - **Discrete GPU:** NVIDIA GeForce RTX 3080 Mobile / Max-Q (8GB/16GB)
  - **RAM:** 32GB
  - **Disk:** 1TB SSD
  - **Shell:** zsh
- Using Arch Linux comes with a cost, the CPU pilots might not be very efficient, stable or up to date.

## Prime number checker



### Experiment Setup

We measured the execution time and derived performance index using:

- **Static load balancing (block-wise)**
- **Static load balancing (cyclic with task size 10)**
- **Dynamic load balancing (task size 10)**

Each implementation was tested using 1, 2, 4, 6, 8, 10, 12, 14, 16, and 32 threads, with execution times measured in milliseconds (ms).

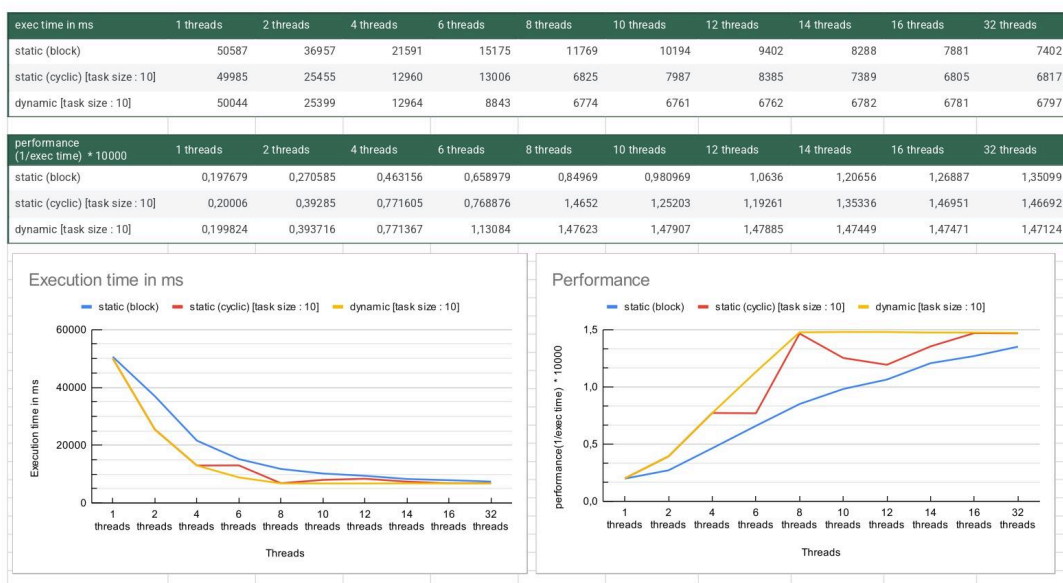
The performance index was defined as:

**- Performance Index = (1/Execution Time) × 10,000**

This means that higher values indicate better performance.

- All the tests were done with the same number of threads and the same end number.
- The end number was set to **1000000**.
- The performance (1 / exec time) was multiplied by **10000** to have a better view of the evolution of the performance.

## Performance Chart



## Interpretation

- **Static block** shows the worst scalability: the fixed partitioning results in some threads finishing early, causing CPU underutilization.
- **Static cyclic (task size = 10)** improves balance by distributing work in round-robin fashion. This reduces idle time and improves throughput.
- **Dynamic scheduling (task-stealing style)** performs the best or ties with static cyclic at almost all thread counts. It efficiently distributes remaining work to idle threads, particularly useful when some rows are more computationally expensive.
- **Beyond 16 threads, there's almost no performance gain** — this suggests we're reaching the limits of physical cores and memory/cache throughput.

## Conclusion

- **The best performance is achieved with dynamic scheduling using 8 to 16 threads.**
- Static cyclic offers nearly the same benefits with slightly more predictable performance curves.
- Static block is easy to implement but inefficient for larger thread counts due to uneven workload distribution.
- The speedup stabilizes after 16 threads due to hardware constraints, indicating that further performance gains would require either different algorithms, GPU acceleration, or architecture-specific optimizations.